

# Senior MLOps Engineer – Take-Home Assessment

## Dataset

[Telco Customer Churn](#)

## Objective

You've just joined a mid-sized telecom company where multiple data scientists are building ML and GenAI models - but without consistent infrastructure. Your job is to start shaping a **minimal MLOps foundation** to help bring order and scalability.

In this assignment, you'll **focus on one core MLOps component** (e.g. a Model Registry or Feature Store) and **mock or describe the others**.

## Scenario

The company needs a way to:

- Register and version models
- Reuse features
- Monitor model health in production
- Retrain and redeploy models automatically

You'll build a **minimal but functional version** of **one** of these components and sketch out how the others would work.

## Your Task

In the attached Jupyter notebook **mlops\_take\_home\_assignment.ipynb** you will find a basic churn prediction model for the given dataset. Your task is to professionalize this setup using MLOps best practices. Pick **ONE** of the following components (**Model Registry or Feature Store**) to implement in code (your choice) and **mock other components**, just include brief placeholder implementations or markdown descriptions to show how you'd approach them in a full setup:

### 1. Model Registry

Use MLflow or your own versioning mechanism

Log: model version, training metrics, parameters, timestamp

Register at least one trained model (any scikit-learn model is fine)

## 2. **Feature Store**

Define and persist a few reusable features

Use lightweight storage (e.g. DuckDB/SQLite) or open-source tools (e.g. Feast)

Provide CLI/script to generate features for both training and inference

## 3. **Monitoring**

Simulate basic monitoring using dummy or replayed inference data

Track: inference volume, latency, simple performance metrics or drift

## 4. **Orchestration and Deployment**

Build a minimal retraining pipeline using Airflow or Python scripts

Retrain model on updated data, register a new version

## 5. **Deployment**

Implement a CI/CD pipeline for containerized inference of the model, using e.g. docker, docker-compose, kubernetes, github actions, gitlab ci/cd...

## **Deliverables**

Submit a GitHub repo or ZIP file with:

- Your code (for the one component you implemented)
- A README.md with:
  - Setup instructions
  - Which component you implemented and why
  - How to run it
  - Short descriptions or pseudocode for the other components

## **Time Estimate**

**~2 hours.** Focus on writing **clean, maintainable code** for one component. No need for perfect coverage or polishing—this is about demonstrating your architectural thinking and hands-on skills.